

# EECS 498-008 / 598-008

# Deep Learning for Computer Vision

## Lecture 1: Introduction

# Logistics:

- Remote for first two weeks (Lectures 1 – 3)
- After that, in-person lecture in Chrysler 220

# Deep Learning for Computer Vision

# Deep Learning for Computer Vision

Building artificial systems  
that process, perceive, and  
reason about visual data


# Computer Vision is everywhere!



Left to right:  
[Image by Roger H Goun](#) is licensed under CC BY 2.0  
[Image](#) is CCO 1.0 public domain  
[Image](#) is CCO 1.0 public domain  
[Image](#) is CCO 1.0 public domain



Left to right:  
[Image](#) is free to use  
[Image](#) is CCO 1.0 public domain  
[Image by NASA](#) is licensed under CC BY 2.0  
[Image](#) is CCO 1.0 public domain



Bottom row, left to right:  
[Image](#) is CCO 1.0 public domain  
[Image by Derek Keats](#) is licensed under CC BY 2.0; changes made  
[Image](#) is public domain  
[Image](#) is licensed under CC-BY 2.0; changes made

# Deep Learning for Computer Vision

Building artificial systems that  
learn from data and experience

# Deep Learning for Computer Vision

Hierarchical learning algorithms  
with many “layers”, (very) loosely  
inspired by the brain

# Artificial Intelligence

Artificial Intelligence

Machine Learning

Computer  
Vision

Artificial Intelligence

Computer  
Vision

Machine Learning

Deep  
Learning


Artificial Intelligence


Computer  
Vision

Machine Learning


Deep  
Learning


This class





Artificial Intelligence






# Today's Agenda

- A brief history of computer vision and deep learning
- Course overview and logistics



# Hubel and Wiesel, 1959

Measure  
brain activity




Cat image by CNX OpenStax is licensed under CC BY 4.0; changes made

1959  
Hubel & Wiesel





# Hubel and Wiesel, 1959

Measure  
brain activity




Cat image by CNX OpenStax is licensed under CC BY 4.0; changes made

1959  
Hubel & Wiesel




**Simple cells:**  
Response to light  
orientation


**Complex cells:**  
Response to light  
orientation and  
movement




# Larry Roberts, 1963




(a) Original picture



(b) Differentiated picture



(c) Feature points selected



Lawrence Gilman Roberts, "Machine Perception of Three-Dimensional Solids", 1963

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

PROJECT MAC

Artificial Intelligence Group  
Vision Memo. No. 100.

July 7, 1966

THE SUMMER VISION PROJECT




Seymour Papert

The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

1959  
Hubel & Wiesel


1963  
Roberts

<https://dspace.mit.edu/handle/1721.1/6125>




Stages of Visual Representation, David Marr, 1970s


# Recognition via Parts (1970s)



Generalized Cylinders,  
Brooks and Binford, 1979



Pictorial Structures,  
Fischler and Elshlager, 1973



# Recognition via Edge Detection (1980s)



1959  
Hubel & Wiesel

1963  
Roberts

1970s  
David Marr

1979  
Gen. Cylinders

1986  
Canny

John Canny, 1986  
David Lowe, 1987

Image is CC0 1.0 public domain

# Recognition via Grouping (1990s)



1959  
Hubel & Wiesel

1963  
Roberts

1970s  
David Marr

1979  
Gen. Cylinders

1986  
Canny

1997  
Norm. Cuts

AI Winter

Normalized Cuts, Shi and Malik, 1997

# Recognition via Matching (2000s)



[Image](#) is public domain



[Image](#) is public domain

1959  
Hubel & Wiesel

1963  
Roberts

1970s  
David Marr

1979  
Gen. Cylinders

1986  
Canny

1997  
Norm. Cuts

1999  
SIFT

AI Winter

SIFT, David  
Lowe, 1999

# Face Detection

Viola and Jones, 2001

One of the first successful applications of machine learning to vision



1959  
Hubel & Wiesel

1963  
Roberts

1970s  
David Marr

1979  
Gen. Cylinders

1986  
Canny




1997  
Norm. Cuts

1999  
SIFT

2001  
V&J

AI Winter

# PASCAL Visual Object Challenge




# IMAGENET Large Scale Visual Recognition Challenge

The Image Classification Challenge:  
1,000 object classes  
1,431,167 images





Output:  
Scale  
T-shirt  
Steel drum  
Drumstick  
Mud turtle


Deng et al, 2009  
Russakovsky et al. IJCV 2015




# IMAGENET Large Scale Visual Recognition Challenge



# AlexNet: Deep Learning Goes Mainstream



Krizhevsky, Sutskever, and Hinton, NeurIPS 2012



# Perceptron



One of the earliest algorithms that could learn from data

Implemented in hardware! Weights stored in potentiometers,  
updated with electric motors during learning

Connected to a camera that used 20x20 cadmium sulfide  
photocells to make a 400-pixel image


Could learn to recognize letters of the alphabet

Today we would recognize it as a **linear classifier**





# Minsky and Papert, 1969

X	Y	F(x,y)
0	0	0
0	1	1
1	0	1
1	1	0



Showed that Perceptrons could not learn the XOR function  
Caused a lot of disillusionment in the field





# Neocognitron: Fukushima, 1980



Computational model the visual system,  
directly inspired by Hubel and Wiesel's  
hierarchy of complex and simple cells

Interleaved simple cells (convolution)  
and complex cells (pooling)

No practical training algorithm





# Neocognitron: Fukushima, 1980




# Backprop: Rumelhart, Hinton, and Williams, 1986

Introduced backpropagation for computing gradients in neural networks


Successfully trained perceptrons with multiple layers



# Convolutional Networks: LeCun et al, 1998



Applied backprop algorithm to a Neocognitron-like architecture  
Learned to recognize handwritten digits  
Was deployed in a commercial system by NEC, processed handwritten checks  
Very similar to our modern convolutional networks!





# 2000s: “Deep Learning”

People tried to train neural networks that were deeper and deeper


Not a mainstream research topic at this time

Hinton and Salakhutdinov, 2006  
Bengio et al, 2007  
Lee et al, 2009  
Glorot and Bengio, 2010




# 2012 to Present: Deep Learning Explosion


CVPR Papers



ML+AI arXiv papers per month



Publications at top Computer Vision conference




# 2012 to Present: Deep Learning is Everywhere

Image Classification




Image Retrieval



Figures copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.


# 2012 to Present: Deep Learning is Everywhere

Object Detection



Ren, He, Girshick, and Sun, 2015


Image Segmentation



Fabaret et al, 2012


# 2012 to Present: Deep Learning is Everywhere

Video Classification



Simonyan et al, 2014

Activity Recognition




# 2012 to Present: Deep Learning is Everywhere

Pose Recognition (Toshev and Szegedy, 2014)




Playing Atari games (Guo et al, 2014)



# 2012 to Present: Deep Learning is Everywhere

Medical Imaging



Levy et al, 2016

Figure reproduced with permission

Whale recognition



[Kaggle Challenge](#)

This image by Christin Khan is in the public domain and originally came from the U.S. NOAA.

Dieleman et al, 2014

From left to right: [public domain by NASA](#), [usage permitted by ESA/Hubble](#), [public domain by NASA](#), and [public domain](#).



# 2012 to Present: Deep Learning is Everywhere



*A white teddy bear  
sitting in the grass*



*A man in a baseball  
uniform throwing a ball*




*A woman is holding  
a cat in her hand*



*A man riding a wave  
on top of a surfboard*



*A cat sitting on a  
suitcase on the floor*



*A woman standing on a  
beach holding a surfboard*

## Image Captioning


Vinyals et al, 2015

Karpathy and Fei-Fei, 2015


All images are CC0 Public domain:

<https://pixabay.com/en/luggage-antique-cat-1643010/>  
<https://pixabay.com/en/teddy-plush-bears-cute-teddy-bear-1623436/>  
<https://pixabay.com/en/surf-wave-summer-sport-litoral-1668716/>  
<https://pixabay.com/en/woman-female-model-portrait-adult-983967/>  
<https://pixabay.com/en/handstand-lake-meditation-496008/>  
<https://pixabay.com/en/baseball-player-shortstop-infield-1045263/>

Captions generated by Justin Johnson using [Neuraltalk2](#)



Figures copyright Justin Johnson, 2015. Reproduced with permission. Generated using the Inceptionism approach from a [blog post](#) by Google Research.




Original image is CC0 public domain  
Starry Night and Tree Roots by Van Gogh are in the public domain  
Bokeh image is in the public domain  
Stylized images copyright Justin Johnson, 2017;  
reproduced with permission



Mordvinsev et al, 2015  
Gatys et al, 2016

# 2012 to Present: Deep Learning is Everywhere



Karras et al, "Progressive Growing of GANs for Improved Quality, Stability, and Variation", ICLR 2018

# 2012 to Present: Deep Learning is Everywhere

## TEXT PROMPT

an armchair in the shape of an avocado. an armchair imitating an avocado.

## AI-GENERATED IMAGES



Ramesh et al, "DALL-E: Creating Images from Text", 2021. <https://openai.com/blog/dall-e/>

# 2012 to Present: Deep Learning is Everywhere

## TEXT PROMPT


an armchair in the shape of a peach. an armchair imitating a peach.

## AI-GENERATED IMAGES




Ramesh et al, "DALL-E: Creating Images from Text", 2021. <https://openai.com/blog/dall-e/>

# Algorithms



Data




Computation




# GFLOP per Dollar

● CPU   ● GPU (FP32)



# GFLOP per Dollar

● CPU    ● GPU (FP32)    ● GPU (Tensor Core)



Recent GPUs have  
“Tensor Cores”:  
Special hardware  
for deep learning!

# 2018 Turing Award




Yoshua Bengio



Geoffrey Hinton




Yann LeCun



Despite our success, computer vision still has a long way to go...

# Computer Vision can cause harm

## Harmful Stereotypes



Barocas et al, "The Problem With Bias: Allocative Versus Representational Harms in Machine Learning", SIGCIS 2017

Kate Crawford, "The Trouble with Bias", NeurIPS 2017 Keynote


Source: <https://twitter.com/jackyalcine/status/615329515909156865> (2015)

## Affect people's lives

### Technology

## A face-scanning algorithm increasingly decides whether you deserve the job

HireVue claims it uses artificial intelligence to decide who's best for a job. Outside experts call it 'profoundly disturbing.'



Source: <https://www.washingtonpost.com/technology/2019/10/22/ai-hiring-face-scanning-algorithm-increasingly-decides-whether-you-deserve-job/>

<https://www.hirevue.com/platform/online-video-interviewing-software>

Example Credit: Timnit Gebru




This image is copyright-free United States government work

Example credit:  
Andrej Karpathy




# Today's Agenda

- A brief history of computer vision and deep learning
- Course overview and logistics



# Today's Agenda

- A brief history of computer vision and deep learning
- Course overview and logistics



# Course Staff

## GSIs / IAs

### Instructor



Justin Johnson  
Assistant Professor, CSE



Karan Desai (KD)



Janpreet Singh (JS)



Jim Yang



Wallace Sui (WS)



Gaurav Kaul

# How to contact us

- Course Website: <https://web.eecs.umich.edu/~justincj/teaching/eecs498/>
  - Syllabus, schedule, assignments, slides, lecture videos, etc
- Piazza: <https://piazza.com/class/kxtai72amx34p0>
  - (Almost) all questions about the course should go here!
  - We will also use Piazza to communicate with you
  - Use private questions if you want to post code
- EECS Autograder: <https://autograder.io/web/course/151>
  - For turning in homework assignments
- [Google Calendar](#): For office hours (starting next week)
- Email: Only for sensitive, confidential issues

# Course Website: Check the Schedule!



EECS 498.008 / 598.008  
Deep Learning for Computer Vision  
Winter 2022

## Schedule

Lectures are Mondays and Wednesdays, 4:30pm to 6pm. Attendance is not required. Recordings will be posted after each lecture in case you are unable to attend the scheduled time.

Some lectures have reading drawn from the course notes of [Stanford CS 231n](#), written by Andrej Karpathy.

Some lectures have optional reading from the book *Deep Learning* by Ian Goodfellow, Yoshua Bengio, and Aaron Courville (GBC for short). The entire text of the book is [available for free online](#) so you don't need to buy a copy.

Event	Date	Description	Course Materials
Lecture 1	Wednesday January 5	<b>Course Introduction</b> Computer vision overview Historical context Course logistics	<a href="#">[slides]</a> <a href="#">[FA2019 video (public)]</a> <a href="#">[Python tutorial]</a> <a href="#">[GBC Sec 1.2]</a> <a href="#">[GBC Sec 6.6]</a>
Lecture 2	Monday January 10	<b>Image Classification</b> Data-driven approach K-Nearest Neighbor Hyperparameters Cross-validation	<a href="#">[slides]</a> <a href="#">[FA2019 video (public)]</a> <a href="#">[231n Image Classification]</a>
Lecture 3	Wednesday January 12	<b>Linear Classifiers</b> Algebraic / Visual / Geometric viewpoints Softmax / SVM classifiers	<a href="#">[slides]</a> <a href="#">[FA2019 video (public)]</a> <a href="#">[231n Linear Classification]</a>
A1 Due	Friday January 14	<b>Assignment 1 Due</b> PyTorch warmup kNN Classifier	<a href="#">[Assignment 1]</a>
	Monday January 17	<b>No class</b> MLK Day	

<https://web.eecs.umich.edu/~justincj/teaching/eecs498/WI2022/schedule.html>

# Piazza Etiquette


- Post only short snippets of code (< 20-30 lines)
- Ask a specific, concrete question
- Explain what you have tried so far, and what happened
- See StackOverflow guide on asking good questions:  
<https://stackoverflow.com/help/how-to-ask>

# Piazza Etiquette

- Post only short snippets of code (< 20-30 lines)
- Ask a specific, concrete question
- Explain what you have tried so far, and what happened
- See StackOverflow guide on asking good questions:  
<https://stackoverflow.com/help/how-to-ask>
- Don't expect an answer within 30 minutes of posting
- Monday – Friday, 10am – 6pm EST we'll try to answer within 2 hours
- Other times, we'll try to answer within 12 hours

# Optional Textbook

- [\*Deep Learning\*](#) by Goodfellow, Bengio, and Courville
- [Free online](#)



# Course Content and Grading

- 6 programming assignments (10% each)
- Midterm Exam (20%)
- Mini-Project (20%)
- Late policy
  - 3 free late days to use on assignments
  - Once free late days are exhausted, 25% penalty per day

# Programming Assignments

- Python, PyTorch, will use Google Colab
- “Earn your wings” – implement things “from scratch” in early assignments, then use PyTorch in later assignments
- “Challenge Questions”
  - Go above and beyond the basic expectations of each assignment
  - Much higher time/points ratio than other parts of the assignment
  - Not necessary to get an A: will be 5% or less of each assignment

# Midterm Exam

- Written exam testing basic concepts from first half of course
- True / False, Multiple choice, short answer
- We will provide a practice exam to get a sense of the format of the questions (but not necessarily the length or difficulty)

# Mini-Project

- New this year!
- Work in teams of 1 to 3
- Homework assignments: We provide starter code, you “fill in the blank”
- Mini-Project: We provide high-level written description of an algorithm to implement, you do the whole thing “from scratch”
- We will give ~3 project descriptions, you pick one of those
- Deliverable: zip of code, and Colab notebook that walks through your implementation and main results (should be like notebooks from homework!)

# Collaboration Policy


- **Rule 1:** Don't look at solutions or code that are not your own; everything you submit should be your own work
- **Rule 2:** Don't share your solution code with others; however discussing ideas or general strategies is fine and encouraged
- **Rule 3:** Indicate in your submissions anyone you worked with
- Turning in something late / incomplete is better than violating the honor code

# Course Philosophy

- Thorough and Detailed.
  - This not “Learn PyTorch in 90 days”, nor “Deep Learning in 10 lines of code”
  - Understand how to write from scratch, debug, and train convolutional and other types of deep neural networks
  - We prefer to write from scratch, rather than rely on existing implementations
- Practical
  - Focus on practical techniques for training and debugging neural networks
  - Will use state-of-the-art software tools like PyTorch and TensorFlow
- State of the art
  - Most material we cover is research published in the last 5 years

# Course Philosophy

- Will also cover some fun topics:
  - Image captioning
  - DeepDream, Artistic Style Transfer



# Course Structure

- First half: Fundamentals
  - Details of how to implement and train different types of networks
  - Fully-connected networks, convolutional networks, recurrent networks
  - How to train and debug, very detailed
- Second half: Applications and “Researchy” topics
  - Object detection, image segmentation, 3D vision, videos
  - Attention, Transformers
  - Vision and Language
  - Generative models: GANs, VAEs, etc
  - Less detailed: provide overview and references, but skip some details

# New Topics since FA2020

- Modern CNN architectures
  - SENets, MobileNets, NAS
  - EfficientNets, RegNets, NFNets
- Vision Transformers
  - Architectures: ViT, DeiT, Swin, MViT
  - Applications: DETR
  - MLP-like architectures
- Vision + Language
  - Language-based pretraining: CLIP, ALIGN
- Self-Supervised Learning
  - Contrastive learning
  - Masked autoencoding

# First homework assignment

- Will be released by today or tomorrow
- Due Friday 1/14/2022
- Next lecture will be enough to complete it

# Next time: Image Classification